


```
In [ ]: import numpy as np
        from enum import Enum
        import copy
```

Consider a standard grid world, where only 4 (up, down, left, right) actions are allowed and the agent deterministically moves accordingly, represented as below. Here yellow is the start state and white is the goal state.

Say, we define our MDP as:

- S: 121 (11 x 11) cells
- A: 4 actions (up, down, left, right)
- P: Deterministic transition probability
- R: -1 at every step
- gamma: 0.9

Our goal is to find an optimal policy (shown in right).



```
In [ ]: # Above grid is defined as below:
        #   - 0 denotes an navigable tile
        #   - 1 denotes an obstruction/wall
        #   - 2 denotes the start state
        #   - 3 denotes an goal state

        # Note: Here the upper left corner is defined as (0, 0)
        #         and lower right corner as (m-1, n-1)

        # Optimal Path: RIGHT RIGHT UP UP LEFT LEFT UP UP UP UP UP UP LEFT LEFT DOWN DOWN LEFT

        GRID_WORLD = np.array([
            [1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1],
            [1, 0, 0, 0, 0, 0, 0, 0, 1, 0, 1],
            [1, 1, 1, 0, 1, 0, 1, 1, 1, 0, 1],
            [1, 3, 0, 0, 1, 0, 1, 0, 1, 0, 1],
```

```

[1, 1, 1, 1, 1, 0, 1, 0, 1, 0, 1],
[1, 0, 0, 0, 1, 0, 1, 0, 0, 0, 1],
[1, 1, 1, 0, 1, 0, 1, 0, 1, 1, 1],
[1, 0, 0, 0, 0, 0, 0, 0, 1, 0, 1],
[1, 1, 1, 1, 1, 1, 1, 0, 1, 0, 1],
[1, 0, 0, 0, 0, 2, 0, 0, 0, 0, 1],
[1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1]
)

```

Actions

```

In [ ]: class Actions(Enum):
    UP    = (0, (-1, 0)) # index = 0, (yaxis_move = -1 and xaxis_move = 0)
    DOWN  = (1, (1, 0))  # index = 1, (yaxis_move = 1 and xaxis_move = 0)
    LEFT  = (2, (0, -1)) # index = 2, (yaxis_move = 0 and xaxis_move = -1)
    RIGHT = (3, (0, 1))  # index = 3, (yaxis_move = 0 and xaxis_move = 1)

    def get_action_dir(self):
        _, direction = self.value
        return direction

    @property
    def index(self):
        indx, _ = self.value
        return indx

    @classmethod
    def from_index(cls, index):
        action_index_map = {a.index: a for a in cls}
        return action_index_map[index]

```

```

In [ ]: # How to use Action enum
for a in Actions:
    print(f"name: {a.name}, action_id: {a.index}, direction_to_move: {a.get_action_dir()}")

print("\n-----\n")

# find action enum from index 0
a = Actions.from_index(0)
print(f"0 index action is: {a.name}")

name: UP, action_id: 0, direction_to_move: (-1, 0)
name: DOWN, action_id: 1, direction_to_move: (1, 0)
name: LEFT, action_id: 2, direction_to_move: (0, -1)
name: RIGHT, action_id: 3, direction_to_move: (0, 1)

-----

0 index action is: UP

```

Policy

```

In [ ]: class BasePolicy:
    def update(self, *args):
        pass

    def select_action(self, state_id: int) -> int:
        raise NotImplemented

```

```

class DeterministicPolicy(BasePolicy):
    def __init__(self, actions: np.ndarray):
        # actions: its a 1d array (/S/ size) which contains action for each state
        self.actions = actions

    def update(self, state_id, action_id):
        assert state_id < len(self.actions), f"Invalid state_id {state_id}"
        assert action_id < len(Actions), f"Invalid action_id {action_id}"
        self.actions[state_id] = action_id

    def select_action(self, state_id: int) -> int:
        assert state_id < len(self.actions), f"Invalid state_id {state_id}"
        return self.actions[state_id]

```

Environment

```

In [ ]: class Environment:
    def __init__(self, grid):
        self.grid = grid
        m, n = grid.shape
        self.num_states = m*n

    def xy_to_posid(self, x: int, y: int):
        _, n = self.grid.shape
        return x*n + y

    def posid_to_xy(self, posid: int):
        _, n = self.grid.shape
        return (posid // n, posid % n)

    def isvalid_move(self, x: int, y: int):
        m, n = self.grid.shape
        return (x >= 0) and (y >= 0) and (x < m) and (y < n) and (self.grid[x, y] != 1)

    def find_start_xy(self) -> int:
        m, n = self.grid.shape
        for x in range(m):
            for y in range(n):
                if self.grid[x, y] == 2:
                    return (x, y)
        raise Exception("Start position not found.")

    def find_path(self, policy: BasePolicy) -> str:
        max_steps = 50
        steps = 0

        P, R = self.get_transition_prob_and_expected_reward()
        num_actions, num_states = R.shape
        all_possible_state_posids = np.arange(num_states)

        path = ""
        curr_x, curr_y = self.find_start_xy()
        while (self.grid[curr_x, curr_y] != 3) and (steps < max_steps):
            curr_posid = self.xy_to_posid(curr_x, curr_y)
            action_id = policy.select_action(curr_posid)
            next_posid = np.random.choice(
                all_possible_state_posids, p=P[action_id, curr_posid])
            action = Actions.from_index(action_id)
            path += f" {action.name}"

```

```

    curr_x, curr_y = self.posid_to_xy(next_posid)
    steps += 1
    return path

def get_transition_prob_and_expected_reward(self): #  $P(s_{next} | s, a)$ ,  $R(s, a)$ 
    m, n = self.grid.shape
    num_states = m*n
    num_actions = len(Actions)
    P = np.zeros((num_actions, num_states, num_states))
    R = np.zeros((num_actions, num_states))
    for a in Actions:
        for x in range(m):
            for y in range(n):
                xmove_dir, ymove_dir = a.get_action_dir()
                xnew, ynew = x + xmove_dir, y + ymove_dir # find the new co-ordinate after t

                posid = self.xy_to_posid(x, y)
                new_posid = self.xy_to_posid(xnew, ynew)

                if self.grid[x, y] == 3:
                    # the current state is a goal state
                    P[a.index, posid, posid] = 1
                    R[a.index, posid] = 0
                elif (self.grid[x, y] == 1) or (not self.isvalid_move(xnew, ynew)):
                    # the current state is a block state or the next state is invalid
                    P[a.index, posid, posid] = 1
                    R[a.index, posid] = -1
                else:
                    # action a is valid and goes to a new position
                    P[a.index, posid, new_posid] = 1
                    R[a.index, posid] = -1
    return P, R

```

Policy Iteration

```

In [ ]: def policy_evaluation(P: np.ndarray, R: np.ndarray, gamma: float,
                             policy: BasePolicy, theta: float,
                             init_V: np.ndarray=None):
    num_actions, num_states = R.shape

    # Please try different starting point for V you will find it will always
    # converge to the same V_pi value.
    if init_V is None:
        init_V = np.zeros(num_states)
    V = copy.deepcopy(init_V)

    delta = 100.0
    while delta > theta:
        delta = 0.0
        for state_id in range(num_states):
            action_id = policy.select_action(state_id)
            v_old = V[state_id]
            # Following equation is a different way of writing the same equation given in the
            # Note here R is an expected reward term.
            V[state_id] = R[action_id, state_id] + gamma * np.dot(P[action_id, state_id], V)
            delta = max(delta, abs(V[state_id] - v_old))
    return V

```

```

def policy_improvement(P: np.ndarray, R: np.ndarray, gamma: float,
                      policy: BasePolicy, V: np.ndarray):
    num_actions, num_states = R.shape
    policy_stable = True
    for state_id in range(num_states):
        old_action_id = policy.select_action(state_id)

        # your code here
        new_action_id = np.argmax(R[:,state_id] + gamma*np.dot(P[:,state_id],V)) # update n

        policy.update(state_id, new_action_id)
        if old_action_id != new_action_id:
            policy_stable = False
    return policy_stable

def policy_iteration(P: np.ndarray, R: np.ndarray, gamma: float,
                    theta: float=1e-3, init_policy: BasePolicy = None):
    num_actions, num_states = R.shape

    # Please try exploring different policies you will find it will always
    # converge to the same optimal policy for valid states.
    if init_policy is None:
        # Say initial policy = all up actions.
        init_policy = DeterministicPolicy(actions=np.zeros(num_states, dtype=int))

    # creating a copy of a initial policy
    policy = copy.deepcopy(init_policy)
    policy_stable = False
    while not policy_stable:
        V = policy_evaluation(P, R, gamma, policy, theta)
        policy_stable = policy_improvement(P, R, gamma, policy, V)
    return policy, V

```

Value Iteration

```

In [ ]: # Directly find the optimal value function
def get_optimal_value(P: np.ndarray, R: np.ndarray, gamma: float,
                     theta: float, init_V: np.ndarray=None):
    num_actions, num_states = R.shape

    # Please try different starting point for V you will find it will always
    # converge to the same V_star value.
    if init_V is None:
        init_V = np.zeros(num_states)
    V = copy.deepcopy(init_V)

    delta = 100.0
    while delta > theta:
        delta = 0.0
        for state_id in range(num_states):
            v_old = V[state_id]
            q_sa = np.zeros(num_actions)
            for a in Actions:
                q_sa[a.index] = R[a.index, state_id] + gamma * np.dot(P[a.index, state_id], V)
            V[state_id] = np.max(q_sa)
            delta = max(delta, abs(V[state_id] - v_old))
    return V

```

```
def value_iteration(P: np.ndarray, R: np.ndarray, gamma: float,
                   theta: float=1e-3, init_V: np.ndarray=None):
    V_star = get_optimal_value(P, R, gamma, theta, init_V)

    num_actions, num_states = R.shape
    policy = DeterministicPolicy(actions=np.zeros(num_states, dtype=int))
    for state_id in range(num_states):
        # Your code here
        action_id = np.argmax(R[:,state_id] + gamma*np.dot(P[:,state_id],V_star)) # update

        policy.update(state_id, action_id)

    return policy, V_star
```

Experiments

```
In [ ]: def is_same_optimal_value(V1, V2, diff_theta=1e-3):
        diff = np.abs(V1 - V2)
        return np.all(diff < diff_theta)
```

```
In [ ]: seed = 0
        np.random.seed(seed)

        gamma = 0.9
        theta = 1e-5
```

```
In [ ]: env = Environment(GRID_WORLD)
        P, R = env.get_transition_prob_and_expected_reward()
```

Exp 1: Using Policy iteration algorithm find the optimal path from start to goal position

```
In [ ]: # # Start with random choice of init_policy.
        # One such choice could be: init_policy = np.ones(env.num_states, dtype=int)
        # Start with your own choice of init_policy
        init_policy = DeterministicPolicy(actions=np.ones(env.num_states, dtype=int))

        pitr_policy, pitr_V_star = policy_iteration(P, R, gamma, theta=theta, init_policy=init_
        pitr_path = env.find_path(pitr_policy)
        print(pitr_path)
```

RIGHT RIGHT UP UP LEFT LEFT UP UP UP UP UP UP LEFT LEFT DOWN DOWN LEFT LEFT

Exp 2: Using value iteration algorithm find the optimal path from start to goal position

```
In [ ]: vitr_policy, vitr_V_star = value_iteration(P, R, gamma, theta=theta)
        vitr_path = env.find_path(vitr_policy)
        print(vitr_path)
```

RIGHT RIGHT UP UP LEFT LEFT UP UP UP UP UP UP LEFT LEFT DOWN DOWN LEFT LEFT

Exp 3: Compare the optimal value function of policy iteration and value iteration algorithm

```
In [ ]: is_same_optimal_value(pitr_V_star, vitr_V_star)
```

Out[]: True

Exp 4: Using initial guess for V as random values, find the optimal value function using policy evaluation and compare it with the optimal value function

```
In [ ]: # Start with random choice of init_V.
# One such choice could be: init_V = np.random.randn(env.num_states)
# Another choice could be: init_V = 10*np.ones(env.num_states)
# Start with your own choice of init_V
init_V = 10*np.ones(env.num_states) # your choice

V_star = policy_evaluation(P, R, gamma, pitr_policy, theta, init_V)
is_same_optimal_value(pitr_V_star, V_star)
```

Out[]: True

Exp 5: Using initial guess for V as random values, find the optimal value function using get_optimal_value and compare it with the optimal value function

```
In [ ]: # Start with random choice.
# One such choice could be: init_V = np.random.randn(env.num_states)
# Another choice could be: init_V = 10*np.ones(env.num_states)
# Start with your own choice of init_V
init_V = 100*np.ones(env.num_states)

V_star = get_optimal_value(P, R, gamma, theta, init_V)
is_same_optimal_value(vitr_V_star, V_star)
```

Out[]: True

In []:

Exp Optional: Try changing the grid by adding multiple paths to the goal state and check if our policy_iteration or value_iteration algorithm is able to find optimal path. Redo the above experiments.

- 1 way to add another path would be GRID_WORLD[4, 1] = 0

In []: