

Tutorial 5 - Options Intro

Please complete this tutorial to get an overview of options and an implementation of SMDP Q-Learning and Intra-Option Q-Learning.

References:

[Recent Advances in Hierarchical Reinforcement Learning](#) is a strong recommendation for topics in HRL that was covered in class. Watch Prof. Ravi's lectures on moodle or nptel for further understanding the core concepts. Contact the TAs for further resources if needed.

```
In [1]: '''
A bunch of imports, you don't have to worry about these
'''

import numpy as np
import random
import gym
# from gym.wrappers import Monitor
import glob
import io
import matplotlib.pyplot as plt
from IPython.display import HTML
```

```
In [2]: '''
The environment used here is extremely similar to the openai gym ones.
At first glance it might look slightly different.
The usual commands we use for our experiments are added to this cell to a
work using this environment.
'''

#Setting up the environment
from gym.envs.toy_text.cliffwalking import CliffWalkingEnv
env = CliffWalkingEnv()

env.reset()

#Current State
print(env.s)

# 4x12 grid = 48 states
print ("Number of states:", env.nS)

# Primitive Actions
action = ["up", "right", "down", "left"]
#correspond to [0,1,2,3] that's actually passed to the environment

# either go left, up, down or right
print ("Number of actions that an agent can take:", env.nA)

# Example Transitions
rnd_action = random.randint(0, 3)
print ("Action taken:", action[rnd_action])
next_state, reward, _, is_terminal, t_prob = env.step(rnd_action)
```

```

print ("Transition probability:", t_prob)
print ("Next state:", next_state)
print ("Reward recieved:", reward)
print ("Terminal state:", is_terminal)
# env.render()

```

```

36
Number of states: 48
Number of actions that an agent can take: 4
Action taken: left
Transition probability: {'prob': 1.0}
Next state: 36
Reward recieved: -1
Terminal state: False

```

```

/usr/local/lib/python3.9/dist-packages/ipykernel/ipkernel.py:283: DeprecationWarning: `should_run_async` will not call `transform_cell` automatically in the future. Please pass the result to `transformed_cell` argument and any exception that happen during the transform in `preprocessing_exc_tuple` in IPython 7.17 and above.
    and should_run_async(code)

```

Options

We custom define very simple options here. They might not be the logical options for this settings deliberately chosen to visualise the Q Table better.

```

In [3]: # We are defining two more options here
# Option 1 ["Away"] -> Away from Cliff (ie keep going up)
# Option 2 ["Close"] -> Close to Cliff (ie keep going down)

def Away(env, state):

    optdone = False
    optact = 0

    if int(state/12)==0:
        optdone = True

    return [optact, optdone]

def Close(env, state):

    optdone = False
    optact = 2

    if int(state/12)==2 or int(state/12)==3:
        optdone = True

    return [optact, optdone]

'''
Now the new action space will contain
Primitive Actions: ["up", "right", "down", "left"]
Options: ["Away", "Close"]
Total Actions : ["up", "right", "down", "left", "Away", "Close"]
Corresponding to [0,1,2,3,4,5]
'''

```

```
Out[3]: '\nNow the new action space will contain\nPrimitive Actions: ["up", "right", "down", "left"]\nOptions: ["Away","Close"]\nTotal Actions :["up", "right", "down", "left", "Away", "Close"]\nCorresponding to [0,1,2,3,4,5]\n'
```

Task 1

Complete the code cell below

```
In [4]: #Q-Table: (States x Actions) == (env.ns(48) x total actions(6))
q_values_SMDP = np.zeros((48,6))

#Update_Frequency Data structure? Check TODO 4
updates_smdp = np.zeros((48,6))

# TODO: epsilon-greedy action selection function
def egreedy_policy(q_values,state,epsilon):
    rand = np.random.uniform(1,0)
    if rand < epsilon:
        return np.random.choice(range(6))
    else:
        return np.argmax(q_values[state])
```

Task 2

Below is an incomplete code cell with the flow of SMDP Q-Learning. Complete the cell and train the agent using SMDP Q-Learning algorithm. Keep the **final Q-table** and **Update Frequency** table handy (You'll need it in TODO 4)

```
In [5]: ##### SMDP Q-Learning

# Add parameters you might need here
gamma = 0.9
alpha = 0.4

# Iterate over 1000 episodes
for _ in range(1000):
    state = env.reset()
    done = False

    # While episode is not over
    while not done:

        # Choose action
        action = egreedy_policy(q_values_SMDP, state, epsilon=0.1)

        # Checking if primitive action
        if action < 4:
            # Perform regular Q-Learning update for state-action pair
            next_state, reward, done,_, _ = env.step(action)

            q_values_SMDP[state][action] += alpha*(reward + gamma*(q_valu
                                                    - q_values_SMDP[stat

            updates_smdp[state][action] += 1
```

```

state = next_state

# Checking if action chosen is an option
reward_bar = 0
counter = 0

if action == 4: # action => Away option
    init_state = state
    optdone = False
    while (optdone == False):

        # Think about what this function might do?
        optact, optdone = Away(env, state)
        next_state, reward, done, _, _ = env.step(optact)

        # Is this formulation right? What is this term? No, cumul
        #  $r_{bar} = r_{t+1} + \gamma r_{t+2} + \dots + \gamma^{\tau-1} r_{t+\tau}$ 
        reward_bar += reward*(gamma**counter)
        counter +=1
        # Complete SMDP Q-Learning Update
        # Remember SMDP Updates. When & What do you update?

        state = next_state

        q_values_SMDP[init_state][4] += alpha*(reward_bar + gamma*(q_
                                                - q_values_SMDP[init_

        updates_smdp[init_state][4] += 1

if action == 5: # action => Close option
    init_state = state
    optdone = False
    while (optdone == False):

        # Think about what this function might do?
        optact, optdone = Close(env, state)
        next_state, reward, done, _, _ = env.step(optact)

        # Is this formulation right? What is this term?
        reward_bar += reward*(gamma**counter)
        counter +=1

        # Complete SMDP Q-Learning Update
        # Remember SMDP Updates. When & What do you update?

        state = next_state

        q_values_SMDP[init_state][5] += alpha*(reward_bar + gamma*(q_
                                                - q_values_SMDP[init_

        updates_smdp[init_state][5] += 1

```

Task 3

Using the same options and the SMDP code, implement Intra Option Q-Learning (In the code cell below). You *might not* always have to search through options to find the options with similar policies, think about it. Keep the **final Q-table** and **Update Frequency** table handy (You'll need it in TODO 4)

```
In [42]: q_values = np.zeros((48,6))
updates = np.zeros((48,6))
```

```
In [ ]: ##### Intra-Option Q-Learning

# Add parameters you might need here
gamma = 0.9
alpha = 0.4

# Iterate over 1000 episodes
for i in range(1000):
    state = env.reset()
    done = False
    print(i)
    # While episode is not over
    while not done:

        # Choose action
        action = egreedy_policy(q_values, state, epsilon=0.1)

        # Checking if primitive action
        if action < 4:
            # Perform regular Q-Learning update for state-action pair
            next_state, reward, done, _, _ = env.step(action)

            q_values[state,action] += alpha*(reward + gamma*(np.max(q_val
updates[state][action] += 1
            state = next_state

            if action == 0:
                q_values[state,4] += alpha*(reward + gamma*(np.max(q_values
updates[state][4] += 1

            if action == 2:
                q_values[state,5] += alpha*(reward + gamma*(np.max(q_values
updates[state][5] += 1

        # Checking if action chosen is an option

        if action == 4: # action => Away option

            optdone = False
            while (optdone == False):

                # Think about what this function might do?
                optact,optdone = Away(env,state)
                next_state, reward, done,_, _ = env.step(optact)

                q_values[state,optact] += alpha*(reward + gamma*(np.max(q
updates[state,optact] += 1

            if optdone:
                q_values[state,4] += alpha*(reward + gamma*(np.max(q_va
updates[state][4] += 1
            else:
                q_values[state][4] += alpha*(reward + gamma*(q_values[n
updates[state][4] += 1

        if done:
```

```

        break

    state = next_state

    if action == 5: # action => Close option
        optdone = False
        while (optdone == False):

            # Think about what this function might do?
            optact, optdone = Close(env, state)
            next_state, reward, done, _, _ = env.step(optact)

            q_values[state, optact] += alpha*(reward + gamma*(np.max(q_values[state, :]) - q_values[state, optact]))
            updates[state, optact] += 1

            if optdone:
                q_values[state][5] += alpha*(reward + gamma*(np.max(q_values[state, :]) - q_values[state][5]))
                updates[state][5] += 1
            else:
                q_values[state][5] += alpha*(reward + gamma*(q_values[state, 5] - q_values[state][5]))
                updates[state][5] += 1

            state = next_state

        if done:
            break

```

In []:

Task 4

Compare the two Q-Tables and Update Frequencies and provide comments.

```

In [31]: import seaborn as sns
import matplotlib.pyplot as plt

sns.heatmap(updates_smdp)


```

```

/usr/local/lib/python3.9/dist-packages/ipykernel/ipkernel.py:283: DeprecationWarning: `should_run_async` will not call `transform_cell` automatically in the future. Please pass the result to `transformed_cell` argument and any exception that happen during the transform in `preprocessing_exc_tuple` in IPython 7.17 and above.
    and should_run_async(code)

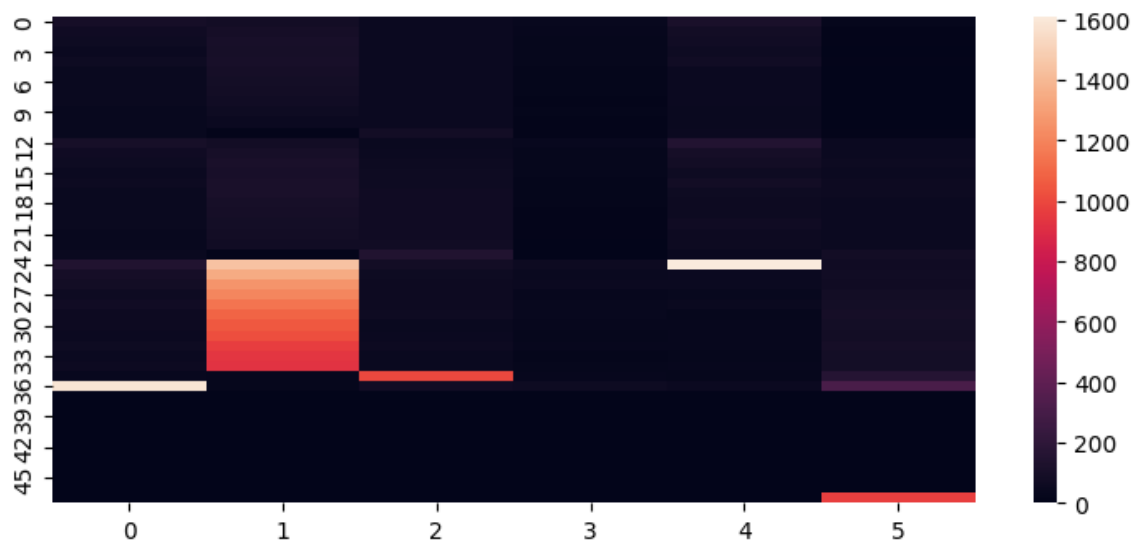
```

Out[31]: <Axes: >



```
In [44]: sns.heatmap(updates_Intra)
```


Out[44]: <Axes: >



```
In [39]: sns.heatmap(q_values_SMDP)
```

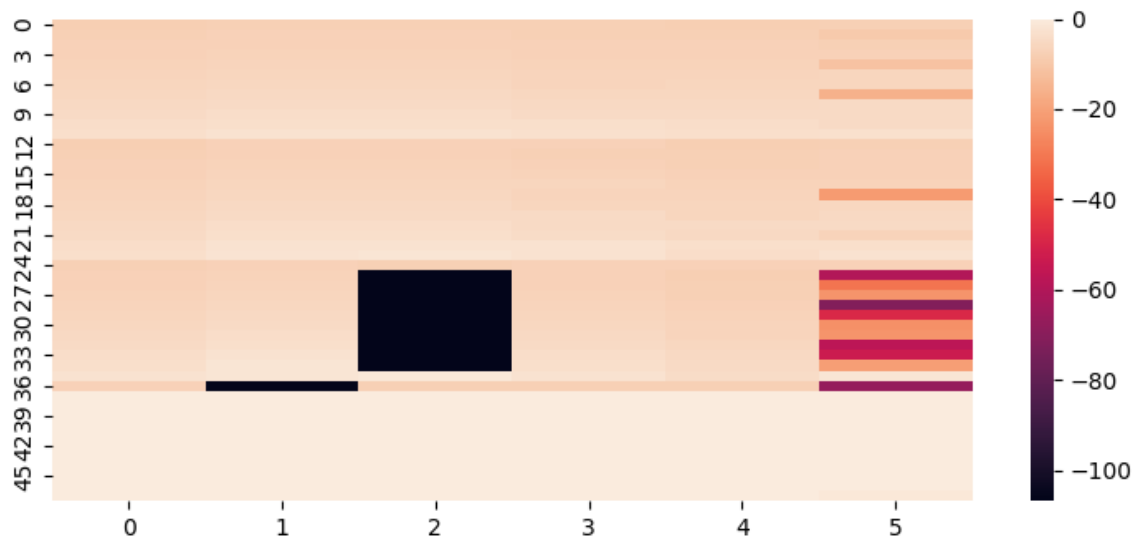
/usr/local/lib/python3.9/dist-packages/ipykernel/ipkernel.py:283: DeprecationWarning: `should_run_async` will not call `transform_cell` automatically in the future. Please pass the result to `transformed_cell` argument and any exception that happen during the transform in `preprocessing_exc_tuple` in IPython 7.17 and above.
and should_run_async(code)

Out[39]: <Axes: >



```
In [45]: sns.heatmap(q_values_Intra)
```

```
Out[45]: <Axes: >
```



Use this text cell for your comments - Task 4

For SMDP we update the q-values after the execution of option with a cumulative reward. We continuously take action over some time steps within the option before updating q-value for that stat-action pair.

In Intra-Option Q-Learning we update q-values for all primitive actions as well as options at each timestep. Due to this the update frequency for Intra-Option Q-Learning is higher than SMDP. But this also depends on the number of times options are chosen over primitive actions.

As for Q-Values since in SMDP they are updated with a cumulative reward after a certain number of time steps, they seem to have a negative value since we get a reward of -1 for each time step and the agent tries to maximize cumulative reward over the subtask or the option. While for Intra-Option learning Q-values seem to be much closer to zero as the agent learns to maximize reward for each time step.


```
In [ ]: !jupyter nbconvert --to html "/content/drive/MyDrive/Colab Notebooks/CS67
```